



Lecture 2: Instruction Set Architecture (ISA) part.1

Jangwoo Kim (Seoul National University)

jangwoo@snu.ac.kr

Slides developed in part by Profs. Austin, Brehob, Falsafi, Hill, Hoe, Lipasti, Martin, Roth, Shen, Smith, Sohi, Tyson, Vijaykumar, and Wenisch @ Carnegie Mellon University, University of Michigan, Purdue University, University of Pennsylvania, University of Wisconsin, POSTECH, and SNU



Office Hour

◆ Office Hour

- Schedule your meeting with TAs and me
- For any matters, we will welcome you

◆ URL

- Use **eTL** actively!
 - Lecture slides will be uploaded
 - Use the Q&A board as much as possible
 - If you have a question, your friends must have the one too.



Announcement

◆ Make-up Lecture (this week only!)

- 3/8 (Thursday 9:30AM) Lecture #2
→ 3/6 (Tuesday 5PM) Lab #1
- 3/6 (Tuesday 5PM) Lab #1
→ 3/8 (Thursday 9:30AM) Lecture #2

Sorry for this sudden notice. :<



[Important] Auditing Policy

◆ Too many "audit" requests!!

- Students from the other session
- Students interested in doing lab only
- Students wanting to study in advance
- Graduate students in other fields

Basic Policy:

(1) If not registered nor audited, DO NOT attend lecture.

**(2) Audit will be allowed only for graduate students
(or some students who I specifically allow)**

(3) Audit students should not take exam nor do lab



Course Plan

◆ 25+ lectures

- Lec .1 ~ 12 :

Microprocessors

(e.g., ISA, CPU structure, Pipelining)

- Lec. 13 ~ 20:

Memory Hierarchy

(e.g., Virtual memory, Cache)

- Lec. 20 ~ end :

Multi-processors

(e.g., Multi-core, cache coherence, GPU)

◆ Lecture + homework

- How modern computer systems work?
- How to design and evaluate a design?

◆ Programming:

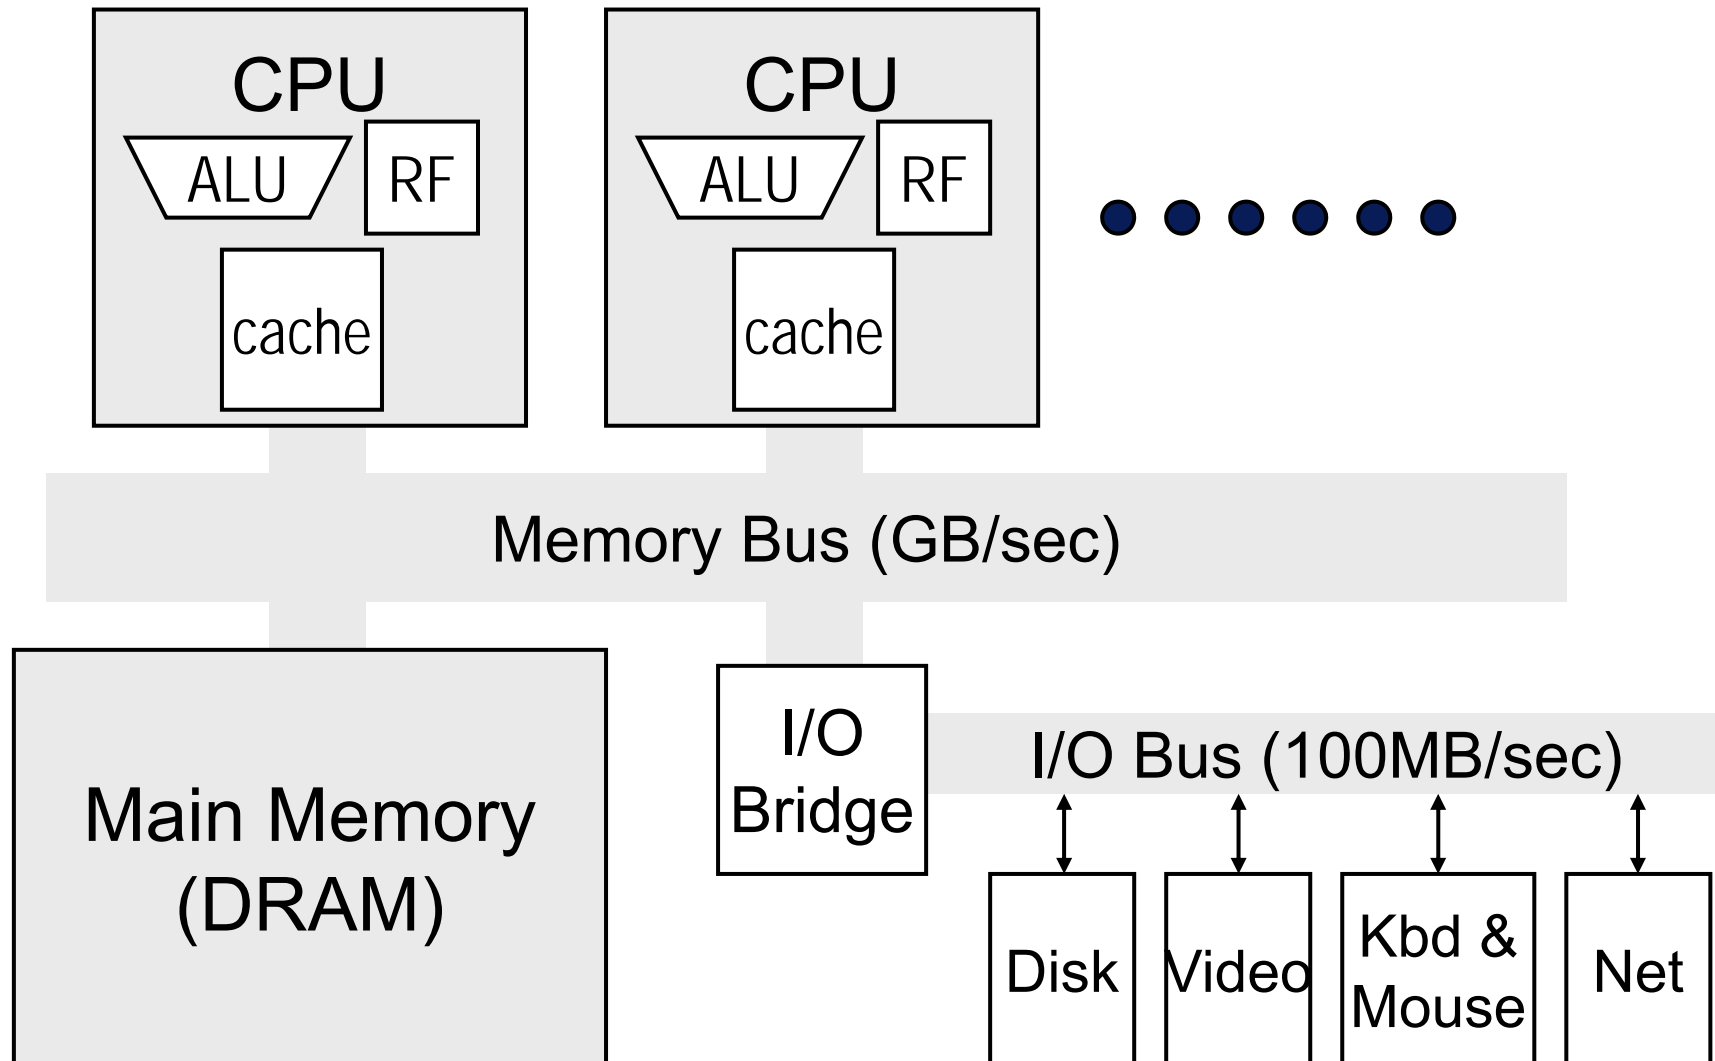
- How to simulate a chip to finalize a design?

◆ Lab:

- How to implement the design on a real chip?



How to specify what a computer does?





How to specify what a computer does?

◆ Architecture Level

- Car: driving manual & operation manual
// you don't have to be a car mechanic to drive a car.
- Computer : ?
// you don't have to be a circuit designer to use a computer.

◆ Microarchitecture (implementation) Level

- A particular car design has a certain configuration of electrical and mechanical components (e.g., v8 engine vs. v4 engine)
- A particular computer design has a certain configuration of datapath and control logic units (e.g., 4M cache vs. 1M cache)

◆ Circuit (=physical implementation) level

- Car: metal sheets, cranks, shafts, gears,
- Computer: gates, wires, semiconductor, transistors, ...



Architecture*

- ◆ “The term *architecture* is used here to describe the attributes of a system as seen by the programmer, i.e., the **conceptual structure** and **functional behavior**, as distinct from the organization of the data flow and controls, the logical design, and the physical implementation.”

--- footnote on page 1, Architecture of the IBM System/360,
Amdahl, Blaauw and Brooks, 1964.



Stored Program (von Neumann) Architecture

- ◆ Stored-program architecture
 - Instructions in a linear memory array
 - Instructions (as well as data) are in memory
 - Instructions can be modified just like data
 - Both forms consist of 0s and 1s

- ◆ Sequential instruction processing
 - Program counter identifies the current instruction
 - Instruction is fetched from memory and executed
 - Program counter is advanced (according to instruction)
 - Repeat



von Neumann vs Dataflow

◆ von Neumann: Consider a von Neumann program

- What is the significance of the program order?
- What is the significance of the storage locations & computations?


"von Neumann bottleneck"

$$\begin{aligned}v &\leftarrow a + b; \\w &\leftarrow b * 2; \\x &\leftarrow v - w \\y &\leftarrow v + w \\z &\leftarrow x * y\end{aligned}$$

◆ Dataflow: Instruction ordering specified by dataflow dependence (no program counter!!)

- Each instruction specifies who should receive results
- An instruction can execute "whenever" all operands are received

Which one is better?





Instruction Set Architecture (ISA)

“User’s manual for the computer”




What are specified/decided in an ISA?

- ◆ Data format and size
 - character, binary, decimal, floating point, negatives
- ◆ “Programmer Visible State” (a.k.a. architectural state)
 - memory, registers, program counter (PC), etc.
- ◆ Instructions: how to “change” the programmer visible state?
 - What to perform and what to perform next
 - Where the operands are
- ◆ Instruction-to-binary (or vice versa) encoding
- ◆ How to interface with the outside world?
- ◆ Protection and privileged operations
- ◆ Software conventions

Very often you compromise immediate optimality
for future scalability and compatibility




Programmer Visible State



Memory

Array of storage locations
indexed by an address



Registers

- Given special names in the ISA
(as opposed to addresses)
- General vs. special purpose

Program Counter (32b~64b)

Memory address
of the current instruction

Instructions (and programs) specify how to transform
the values of programmer visible state




General Instruction Classes

- ◆ Arithmetic and logical operations (e.g., add, sub, and, or)
 - 1) Fetch operands from specified locations
 - 2) Compute a result as a function of the operands
 - 3) Store result to a specified location
 - 4) Update PC to the next sequential instruction
- ◆ Data movement operations (e.g., move, load, store)
 - 1) Fetch operands from specified locations
 - 2) Store operand values to specified locations
 - 3) Update PC to the next sequential instruction
- ◆ Control flow operations (e.g., branch, jump)
 - 1) Fetch operands from specified locations
 - 2) Compute a **branch condition** and a **target address**
 - 3) If “**branch condition** is true” then $PC \leftarrow \text{target address}$
else $PC \leftarrow \text{next seq. instruction}$

Generally defined to be atomic.




An Early ISA: EDSAC (~1950)



◆ Single accumulator architecture, i.e. $ACC \leftarrow ACC \oplus M[n]$

- **A n**: add $M[n]$ into ACC
- **T n**: transfer the contents of ACC to $M[n]$
- **E n**: If $ACC > -1$, branch to $M[n]$ or proceed serially
- **I n**: Read the next character from paper tape, and store it as the least significant 5 bits of $M[n]$
- **Z**: Stop the machine and ring the warning bell




◆ Only “absolute” addressing in data and control transfer

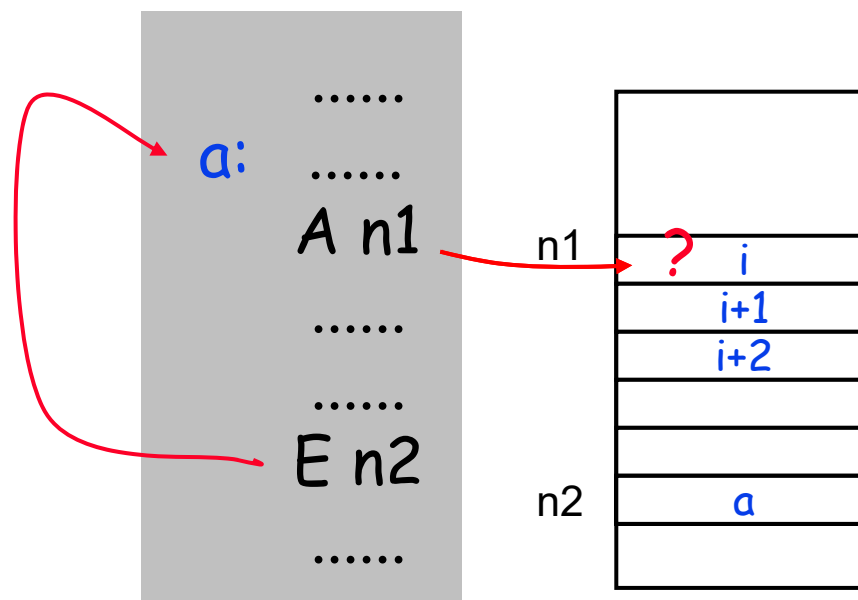


Would you make the same mistake?

◆ Function call



◆ Array access in a loop



How to fix this?



Evolution of “Register” Architecture

- ◆ Accumulator
 - A legacy from the “adding” machine days
“AC” button on your calc?
- ◆ Accumulator + address registers
 - Need **register indirection**
 - Initially address registers were special-purpose,
i.e., can only be loaded with an address for indirection
 - Eventually **arithmetic on addresses** became supported
- ◆ “General purpose registers (GPR)”
 - **All registers good for all purposes**
 - Grew from a few registers to 32 (common for RISC) to 128 (e.g., Intel Itanium)

What are driving the changes?



Operand Sources?

◆ Number of Operands

Monadic	OP in2	(e.g. EDSAC)
Binatic	OP inout, in2	(e.g. IBM 360)
	- 3 operands in a smaller encoding to save memory	
Triadic	OP out, in1, in2	(e.g. MIPS)

◆ Can ALU operands be in memory?

Yes! e.g. x86/VAX/“CISC”

No! e.g. MIPS/“RISC” → “**load-store architecture**”

◆ How many different variations

a very few e.g. MIPS / RISC

a lot e.g. x86

everything goes e.g. VAX



Memory Addressing Modes

- ◆ Absolute `LW rt, 10000` // LW=load word
use immediate value as address
- ◆ Register Indirect: `LW rt, (rbase)`
use $GPR[r_{base}]$ as address
- ◆ Displaced or based: `LW rt, offset(rbase)`
use $offset + GPR[r_{base}]$ as address
- ◆ Indexed: `LW rt, (rbase, rindex)`
use $GPR[r_{base}] + GPR[r_{index}]$ as address
- ◆ Memory Indirect `LW rt ((rbase))`
use value at $M[GPR[r_{base}]]$ as address
- ◆ Auto inc/decrement `LW rt, (rbase);` // $r_{base} ++/--$
use $GPR[r_{base}]$ as address, but inc. or dec. $GPR[r_{base}]$ each time
- ◆ Anything else can you think off



VAX-11: ISA in mid-life crisis

- ◆ First commercial 32-bit machine
considered an important milestone
- ◆ Ultimate in “orthogonality” and “completeness”
All of the above addressing modes x { 7 integer and 2 floating point formats} x {more than 300 opcodes}
- ◆ Opcode in excess
 - 2-operand and 3-operand versions of ALU ops
 - INS(/REM)QUE (for circular doubly-linked list)
 - “polyf”: 4th-degree polynomial solve
- ◆ Encoding
addl3 r1,737(r2),(r3)[r4] 7-byte, sequential decode



MIPS RISC

- ◆ Simple operations
 - 2-input, 1-output arithmetic and logical operations
 - Only few alternatives exist to do the same thing
- ◆ Simple data movements
 - ALU ops are register-to-register (**need a large register file**)
 - Memory can be accessed by only load and store instructions
→ **“Load-store architecture”**
- ◆ Simple branches
 - Limited varieties of branch conditions and targets
- ◆ Simple instruction encoding
 - All instructions encoded **in the same number of bits**
 - Only a few formats

**Such ISA intended for compiler advances
rather than assembly programmers**



Evolution of ISA

- ◆ Why were the earlier ISAs so simple? e.g. EDSAC
 - Technology limitation
 - Inexperience, lack of precedence
- ◆ Why did it get so complicated later? e.g. VAX11
 - Assembly programming
 - Lack of memory size and performance
 - Microprogrammed implementation
- ◆ Why did it become simple again? e.g. RISC
 - Memory size and speed (cache!)
 - **Compilers**
- ◆ Why x86 is still “king of the hill”?
 - Technology vs. economics
 - Technology vs. psychology
 - Technology vs. deep pocket

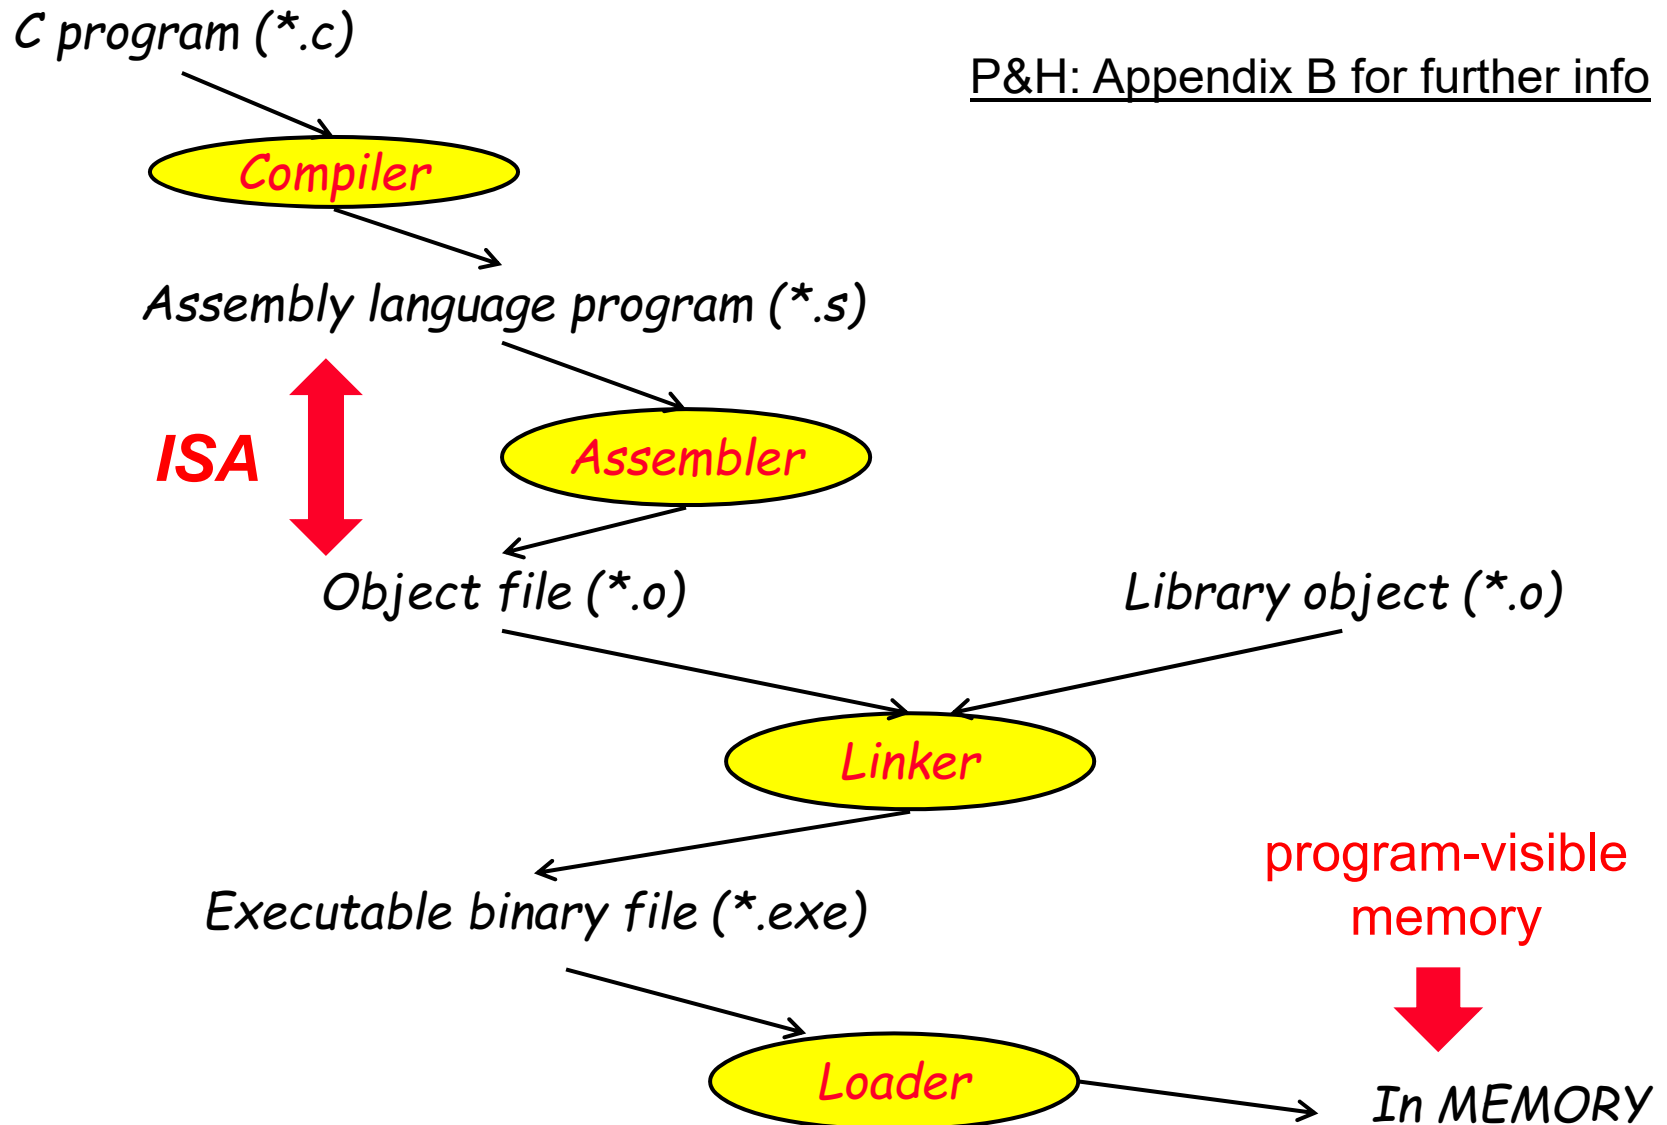
*Complex Instruction
Set Architecture (CISC)*

*Reduced Instruction
Set Architecture (RISC)*



Maybe Review (from programming 101)?

P&H: Appendix B for further info





Example C Program

main.c

```
int buf[2] = {1, 2};

int main()
{
    swap();
    return 0;
}
```

swap.c

```
extern int buf[];

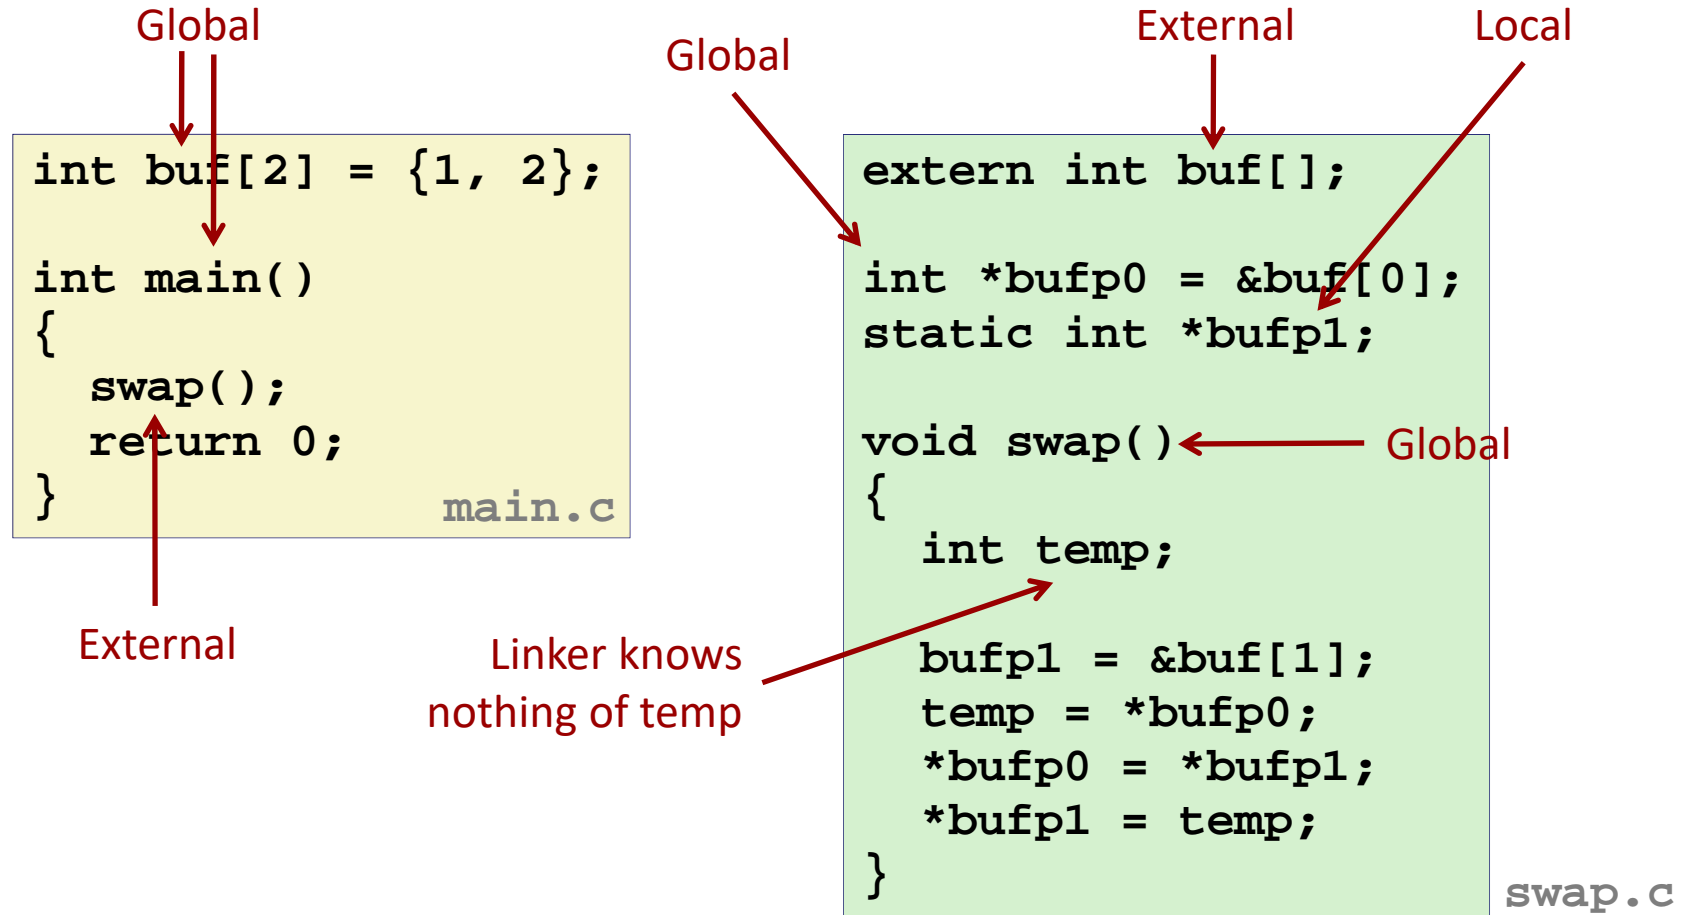
int *bufp0 = &buf[0];
static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```




Resolving Symbols



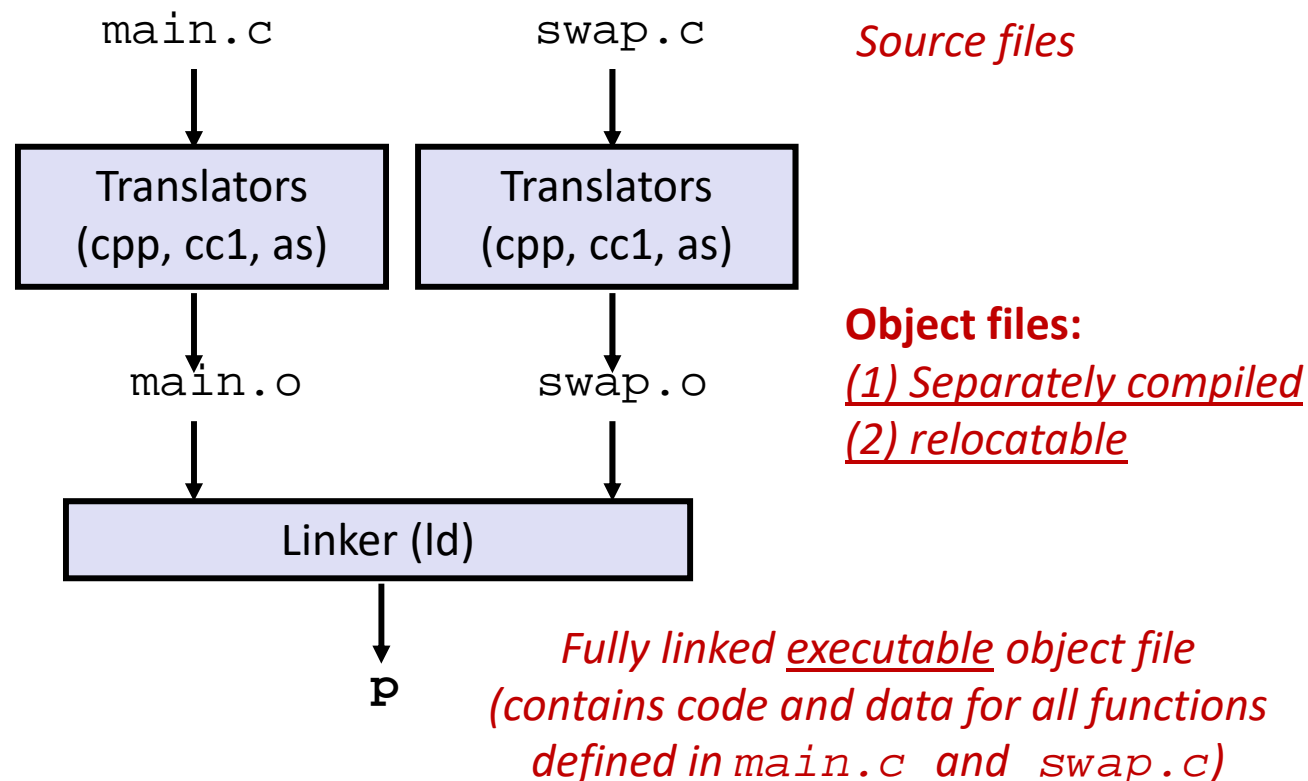
Information necessary to support linking



Static Linking

- ◆ Programs are translated and linked using a *compiler driver*:

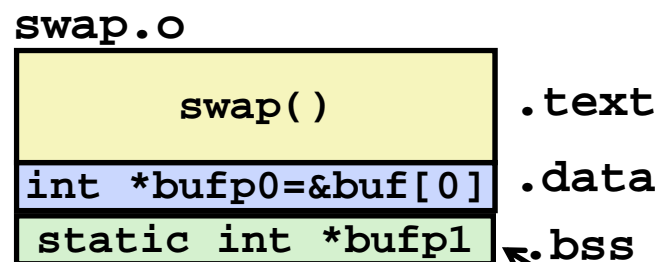


```
- unix> gcc -O2 -g -o p main.c swap.c  
- unix> ./p
```



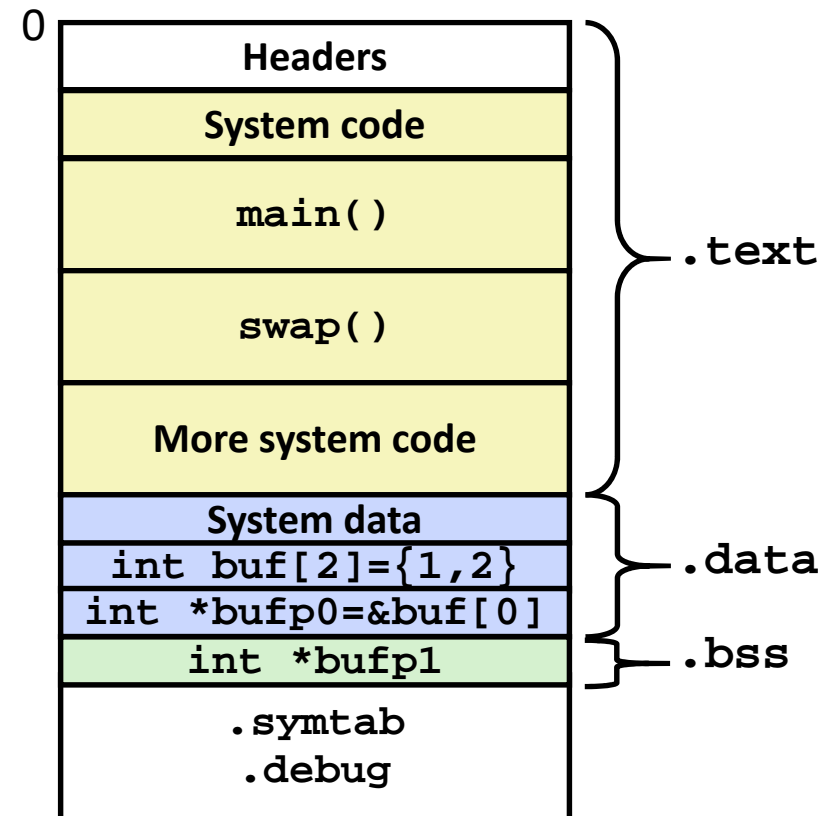


Relocating Code and Data

Relocatable Object Files



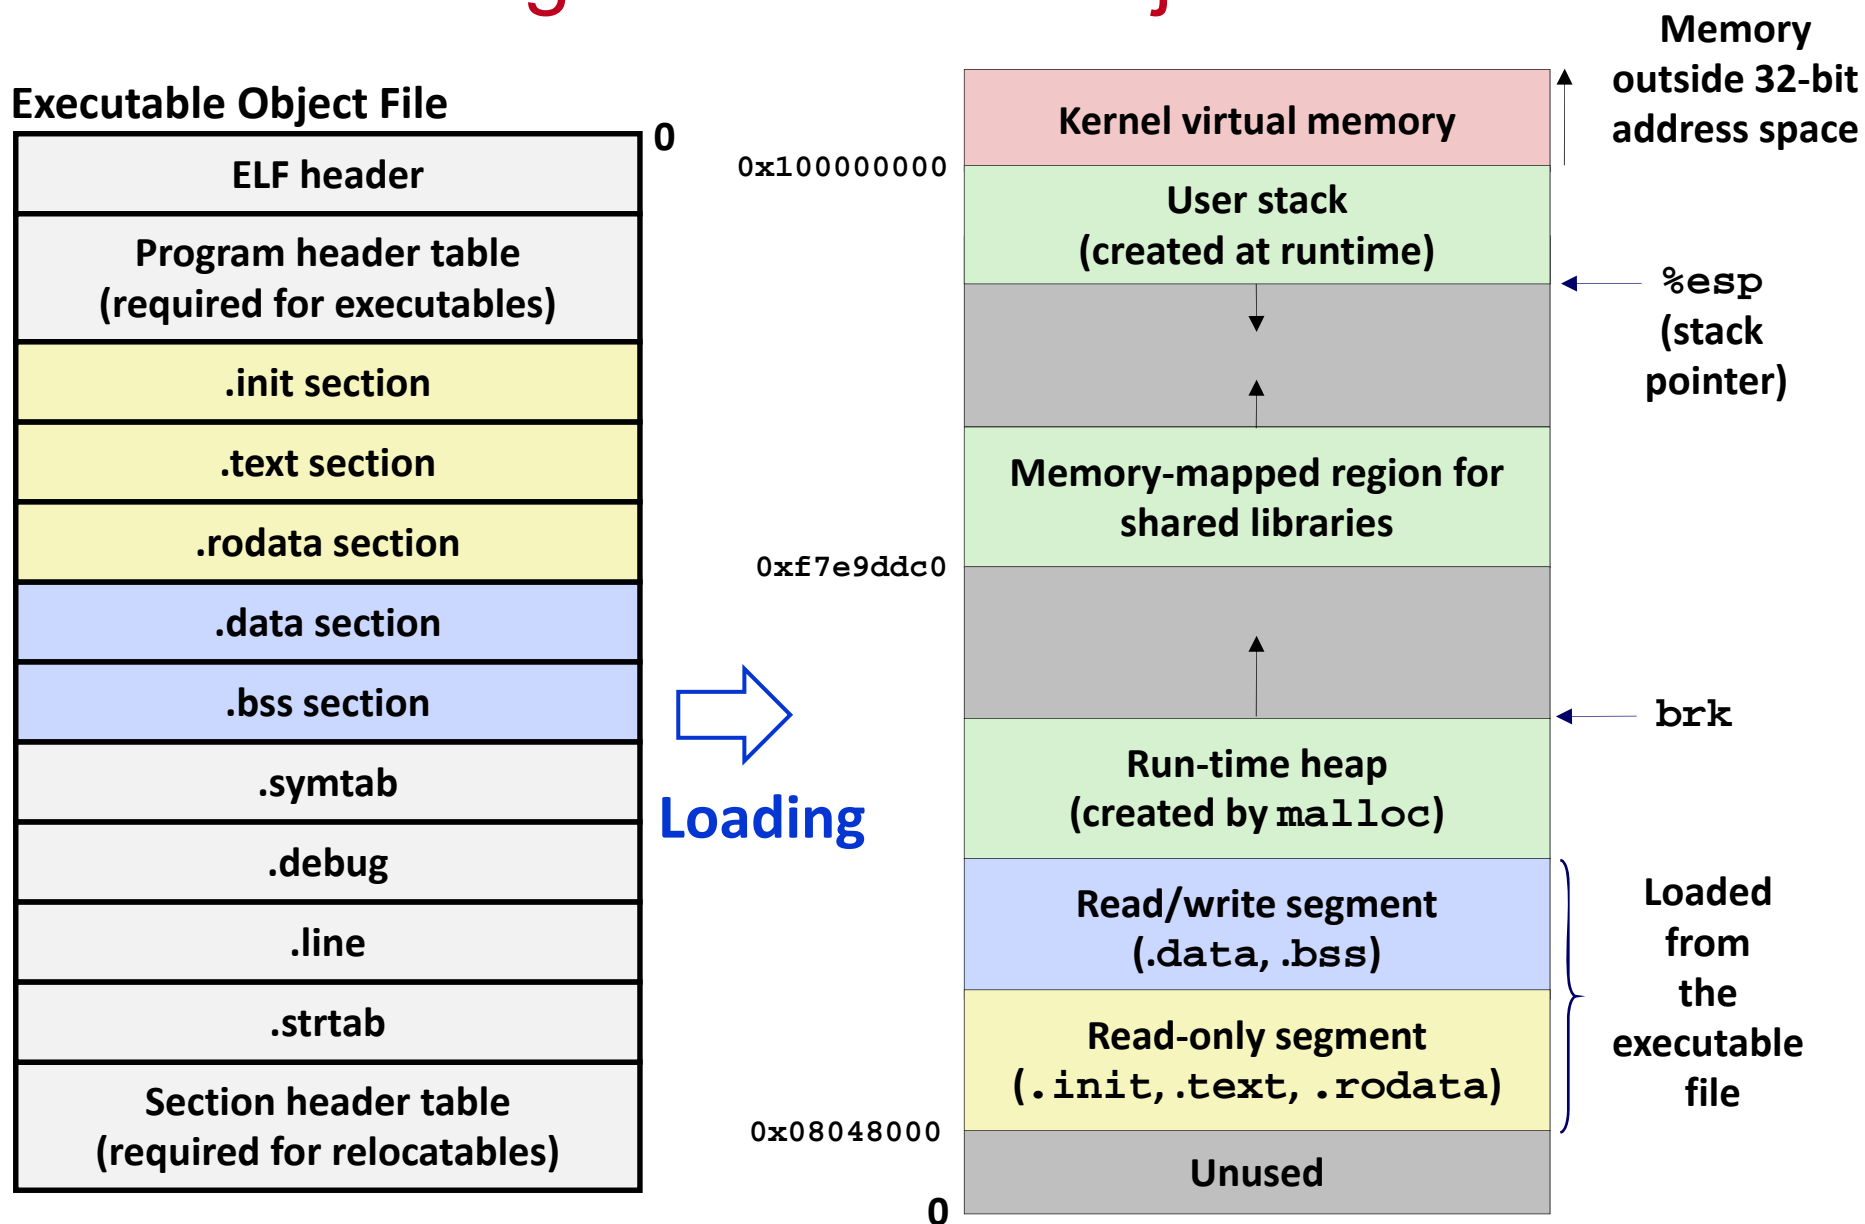
Executable Object File



Even though private to swap, requires allocation in .bss
(block started by symbol for uninitialized static variables)



Loading Executable Object Files





What if you have to design an ISA?



Open-Ended Design: extrapolation and anticipation

“a dependable base for a decade of customer planning and customer programming, and continuing laboratory development...” Typical life expectance 15~20 years, but.....

- ◆ “Asynchronous” operation of components
 - Abstract out exact time, performance etc to allow (1) changing technology and (2) relative speed of components
- ◆ Parameterization of storage capacity, multi CPU, multi I/O, etc
- ◆ Permit future extensions by “reserving” spare bits in instruction encoding
- ◆ Standard interfaces for expansion sub-systems

[Amdahl, Blaauw and Brooks, 1964]



General Purpose

- ◆ General Purpose = effective support for “large and small, separate and mixed applications” in many domains (e.g., commercial, scientific, real-time....)
- ◆ How
 - Code-independent operation
 - No special interpretation of bit pattern in data
e.g. ASCII character has no special significance.
 - Except where essential
e.g., Integer, floating point, etc
 - Support full generality of logic manipulation on bit and data entities
 - Fine-grain memory addressability (down to small units of bits)

[Amdahl, Blaauw and Brooks, 1964]



Inter-Model Compatibility

- ◆ **Strict program compatibility** = “a valid program whose logic will not depend implicitly upon time of execution and which runs upon configuration A, will also run on configuration B if the latter includes at least the required storage, at least the required I/O devices”
- ◆ Invalid programs [...] are not constrained to yield the same result
 - “invalid program” means a program that violates the architecture manual and not that it generates exceptions
 - exceptional conditions are part of the architecture
- ◆ By virtualization of logical structures and functions
- ◆ The King of Binary Compatibility: **Intel x86, IBM 360**
 - software base
 - performance scalability

[Amdahl, Blaauw and Brooks, 1964]



Wrap-up: Terminologies

- ◆ Instruction Set Architecture (ISA)
 - The machine behavior as observable and controllable by the programmer
- ◆ Instruction Set
 - The set of commands understood by the computer
- ◆ Machine Code
 - A collection of instructions encoded in binary format
 - Directly consumable by the hardware
- ◆ Assembly Code
 - A collection of instructions expressed in “textual” format
e.g. Add r1, r2, r3
 - Converted to machine code by an assembler
 - One-to-one correspondence with machine code



Question?

Announcements: ISA part.2 in the next lecture

***Reading: P&H Ch 2 & optionally Appendix E
(summary of all major ISAs)***

Handouts: None