



Lecture 3: Instruction Set Architecture (ISA) part.2 (MIPS)

Jangwoo Kim (Seoul National University)

jangwoo@snu.ac.kr



Announcement

Announcements: HW #1 will be posted on eTL site soon.

*Reading: Start reading P&H Ch.3
(stuff for the next lectures).*

Handouts: None



[Important] Auditing Policy

◆ Too many "audit" requests!!

- Students from the other session
- Students interested in doing lab only
- Students wanting to study in advance
- Graduate students in other fields

Basic Policy:

(1) If not registered nor audited, DO NOT attend lecture.

**(2) Audit will be allowed only for graduate students
(or some students who I specifically allow)**

(3) Audit students should not take exam nor do lab



Review

- ◆ Architecture = programmer visible state
- ◆ What is “programmer visible state?”
 - Who is “programmer” in this context?
 - This is often called as “architectural state”

Useful thinking:

- Writing a system program. What should we know?
 - Operating system, architecture simulator?
- Writing a compiler. What should we know?
 - C code → ?
- Writing an OS for context switching. What should we know?
 - context-out: a program running in CPU → ?
 - context-in: ? → an empty CPU



Instruction Set Architecture

- ◆ A stable platform, typically 15~20 years
 - Guarantees binary compatibility for SW investments
 - Permits adoption of foreseeable technology advances
- ◆ User-level ISA
 - **Program visible state** and instructions available to user processes
 - Single-user abstraction on top of HW/SW virtualization
- ◆ “Virtual Environment” Architecture
 - State and instructions to control virtualization (e.g., caches, sharing)
 - User-level, but not used by your average user programs
- ◆ “Operating Environment” Architecture
 - State and instructions to implement virtualization
 - Privileged/protected access reserved for OS



What are specified/decided in an ISA?

- ◆ Data format and size
 - character, binary, decimal, floating point, negatives
- ◆ “Programmer Visible State”
 - memory, registers, program counters, etc.
- ◆ Instructions: how to change the programmer visible state?
 - what to perform and what to perform next
 - where are the operands
- ◆ Instruction-to-binary encoding
- ◆ How to interface with the outside world?
- ◆ Protection and privileged operations
- ◆ Software conventions

Very often you compromise immediate optimality for future scalability and compatibility



MIPS R2000 Program Visible State

Program Counter

32-bit memory address
of the current instruction

M[0]
M[1]
M[2]
M[3]
M[4]
M[N-1]

****Note**** r0=0

r1

r2

General Purpose
Register File
32 32-bit words
named r0...r31

Memory

2^{32} by 8-bit locations (4 Giga Bytes)
32-bit address



Data Format


- ◆ Most things are 32 bits (= word)
 - Instruction and data addresses
 - Signed and unsigned integers
 - Just bits
- ◆ Also 16-bit word and 8-bit word (= byte)
- ◆ Floating-point numbers
 - IEEE standard 754
 - Single precision:
 - 1-bit sign (S), 8-bit exponent (E), 23-bit fraction (F)
 - $(-1)^S \times F \times 2^{(\text{exponent})}$ ← Not exactly this though.
 - Fraction (aka significand)
 - e.g., $(-1)^1 \times (11) \times 2^{-(10)} = ?$
 - Double precision:
 - 11-bit exponent, 52-bit significand



Big Endian vs. Little Endian

(Part I, Chapter 4, Gulliver's Travels)

- ◆ 32-bit signed or unsigned integer comprises 4 bytes



- ◆ On a byte-addressable machine

Big Endian

Little Endian

MSBAddr→LSB

byte 3	byte 2	byte 1	byte 0
byte 7	byte 6	byte 5	byte 4
byte 11	byte 10	byte 9	byte 8
byte 15	byte 14	byte 13	byte 12
byte 19	byte 18	byte 17	byte 16

pointer points to the big end

LSB	Addr	→	MSB
byte 0	byte 1	byte 2	byte 3
byte 4	byte 5	byte 6	byte 7
byte 8	byte 9	byte 10	byte 11
byte 12	byte 13	byte 14	byte 15
byte 16	byte 17	byte 18	byte 19

pointer points to the little end


What difference does it make? check out `htonl()`, `ntohl()` in `in.h`



Instruction Formats


◆ 3 simple formats

- R-type, 3 register operands




R-type

- I-type, 2 register operands and 16-bit immediate operand



I-type

- J-type, 26-bit immediate operand



J-type

◆ Simple Decoding

- 4 bytes per instruction, regardless of format
- Must be 4-byte aligned (2 lsb of PC must be 2b'00)
- Format and fields readily extractable



ALU Instructions

- ◆ Assembly (e.g., register-register signed addition)

ADD rd_{reg} rs_{reg} rt_{reg}

- ◆ Machine encoding

0	rs	rt	rd	0	ADD
6-bit	5-bit	5-bit	5-bit	5-bit	6-bit

R-type

- ◆ Semantics **// how arch. state is changed**

- $GPR[rd] \leftarrow GPR[rs] + GPR[rt]$
- $PC \leftarrow PC + 4$

- ◆ Exception on “overflow”

- ◆ Variations

- Arithmetic: {signed, unsigned} x {ADD, SUB}
- Logical: {AND, OR, XOR, NOR}
- Shift: {Left, Right-Logical, Right-Arithmetic}



ADD

Add Word

31	26	25	21	20	16	15	11	10	6	5	0
SPECIAL 0 0 0 0 0 0						rs	rt	rd	0 0 0 0 0 0	ADD 1 0 0 0 0 0	
6						5	5	5	5	6	

Format: ADD rd, rs, rt

MIPS I

Purpose: To add 32-bit integers. If overflow occurs, then trap.

Description: $rd \leftarrow rs + rt$

The 32-bit word value in GPR *rt* is added to the 32-bit value in GPR *rs* to produce a 32-bit result. If the addition results in 32-bit 2's complement arithmetic overflow then the destination register is not modified and an Integer Overflow exception occurs. If it does not overflow, the 32-bit result is placed into GPR *rd*.

Restrictions:

On 64-bit processors, if either GPR *rt* or GPR *rs* do not contain sign-extended 32-bit values (bits 63..31 equal), then the result of the operation is undefined.

Operation:

```
if (NotWordValue(GPR[rs]) or NotWordValue(GPR[rt])) then UndefinedResult() endif
temp  $\leftarrow$  GPR[rs] + GPR[rt]
if (32_bit_arithmetic_overflow) then
    SignalException(IntegerOverflow)
else
    GPR[rd]  $\leftarrow$  sign_extend(temp31..0)
endif
```

Exceptions:

Integer Overflow

Programming Notes:

ADDU performs the same arithmetic operation but, does not trap on overflow.

From the
ISA manual



Reg-Reg Instruction Encoding

		SPECIAL function						
		2...0						
5...3	0	1	2	3	4	5	6	7
0	SLL	*	SRL	SRA	SLLV	*	SRLV	SRAV
1	JR	JALR	*	*	SYSCALL	BREAK	*	SYNC
2	MFHI	MTHI	MFLO	MTLO	DSLLV _ε	*	DSRLV _ε	DSRAV _ε
3	MULT	MULTU	DIV	DIVU	DMULT _ε	DMULTU _ε	DDIV _ε	DDIVU _ε
4	ADD	ADDU	SUB	SUBU	AND	OR	XOR	NOR
5	*	*	SLT	SLTU	DADD _ε	DADDU _ε	DSUB _ε	DSUBU _ε
6	TGE	TGEU	TLT	TLTU	TEQ	*	TNE	*
7	DSLL _ε	*	DSRL _ε	DSRA _ε	DSLL32 _ε	*	DSRL32 _ε	DSRA32 _ε

[MIPS R4000 Microprocessor User's Manual]

What patterns do you see? Why are they there?



ALU Instructions

- ◆ Assembly (e.g., regi-immediate signed additions)

ADDI rt_{reg} rs_{reg} $immediate_{16}$

- ◆ Machine encoding

ADDI	rs	rt	immediate
6-bit	5-bit	5-bit	16-bit

I-type

- ◆ Semantics

- $GPR[rt] \leftarrow GPR[rs] + \text{sign-extend}(\text{immediate})$
- $PC \leftarrow PC + 4$

- ◆ Exception on “overflow”

- ◆ Variations

- Arithmetic: {signed, unsigned} x {ADD, ~~SUB~~}
- Logical: {AND, OR, XOR, LUI}



Reg-Immed Instruction Encoding

28...26		Opcode						
31...29	0	1	2	3	4	5	6	7
0	SPECIAL	REGIMM	J	JAL	BEQ	BNE	BLEZ	BGTZ
1	ADDI	ADDIU	SLTI	SLTIU	ANDI	ORI	XORI	LUI
2	COP0	COP1	COP2	*	BEQL	BNEL	BLEZL	BGTZL
3	DADDI ϵ	DADDIU ϵ	LDL ϵ	LDR ϵ	*	*	*	*
4	LB	LH	LWL	LW	LBU	LHU	LWR	LWU ϵ
5	SB	SH	SWL	SW	SDL ϵ	SDR ϵ	SWR	CACHE δ
6	LL	LWC1	LWC2	*	LLD ϵ	LDC1	LDC2	LD ϵ
7	SC	SWC1	SWC2	*	SCD ϵ	SDC1	SDC2	SD ϵ

[MIPS R4000 Microprocessor User's Manual]



Assembly Programming 101

- ◆ Break down high-level program constructs into a sequence of elemental operations

- ◆ E.g. High-level Code

```
f = ( g + h ) - ( i + j )
```

- ◆ Assembly Code

- suppose f, g, h, i, j are in r_f, r_g, r_h, r_i, r_j
- suppose r_{temp} is a free register

```
add r_temp r_g r_h      # r_temp = g+h
add r_f r_i r_j          # r_f = i+j
sub r_f r_temp r_f       # f = r_temp - r_f
```




Load Instructions

- ◆ Assembly (e.g., load 4-byte word)

LW rt_{reg} $offset_{16}$ ($base_{reg}$)

- ◆ Machine encoding

LW	base	rt	offset
6-bit	5-bit	5-bit	16-bit

I-type

- ◆ Semantics

- $effective_address = \text{sign-extend}(offset) + GPR[base]$
- $GPR[rt] \leftarrow MEM[\text{translate}(effective_address)]$
- $PC \leftarrow PC + 4$


- ◆ Exceptions

- Address must be “word-aligned”
What if you want to load an unaligned word?
- MMU exceptions \leftarrow Let's handle unexpected behaviors.



LW

Load Word



Format: LW rt, offset(base)

MIPS I

Purpose: To load a word from memory as a signed value.

Description: $rt \leftarrow \text{memory}[\text{base} + \text{offset}]$

The contents of the 32-bit word at the memory location specified by the aligned effective address are fetched, sign-extended to the GPR register length if necessary, and placed in GPR *rt*. The 16-bit signed *offset* is added to the contents of GPR *base* to form the effective address.

Restrictions:

The effective address must be naturally aligned. If either of the two least-significant bits of the address are non-zero, an Address Error exception occurs.

MIPS IV: The low-order 2 bits of the *offset* field must be zero. If they are not, the result of the instruction is undefined.

Operation: 32-bit processors

$vAddr \leftarrow \text{sign_extend}(\text{offset}) + \text{GPR}[\text{base}]$
 if $(vAddr_{1..0}) \neq 0^2$ then $\text{SignalException}(\text{AddressError})$ endif
 $(pAddr, \text{uncached}) \leftarrow \text{AddressTranslation}(vAddr, \text{DATA}, \text{LOAD})$
 $\text{memword} \leftarrow \text{LoadMemory}(\text{uncached}, \text{WORD}, pAddr, vAddr, \text{DATA})$
 $\text{GPR}[rt] \leftarrow \text{memword}$

Operation: 64-bit processors


$vAddr \leftarrow \text{sign_extend}(\text{offset}) + \text{GPR}[\text{base}]$
 if $(vAddr_{1..0}) \neq 0^2$ then $\text{SignalException}(\text{AddressError})$ endif
 $(pAddr, \text{uncached}) \leftarrow \text{AddressTranslation}(vAddr, \text{DATA}, \text{LOAD})$
 $pAddr \leftarrow pAddr_{\text{PSIZE}-1..3} \parallel (pAddr_{2..0} \text{ xor } (\text{ReverseEndian} \parallel 0^2))$
 $\text{memdouble} \leftarrow \text{LoadMemory}(\text{uncached}, \text{WORD}, pAddr, vAddr, \text{DATA})$
 $\text{byte} \leftarrow vAddr_{2..0} \text{ xor } (\text{BigEndianCPU} \parallel 0^2)$
 $\text{GPR}[rt] \leftarrow \text{sign_extend}(\text{memdouble}_{31+8*\text{byte}..8*\text{byte}})$

Exceptions:


TLB Refill, TLB Invalid
 Bus Error
 Address Error



Data Alignment



- ◆ LW/SW alignment restriction
 - Not optimized to fetch memory bytes not within a word boundary
 - Not optimized to rotate unaligned bytes into registers
- ◆ Can use separate opcodes for the infrequent case



- LWL/LWR is slower but it is okay
- Note LWL and LWR still fetch within word boundary



Store Instructions

- ◆ Assembly (e.g., store 4-byte word)

$SW\ rt_{reg}\ offset_{16}\ (base_{reg})$

- ◆ Machine encoding

SW	base	rt	offset
6-bit	5-bit	5-bit	16-bit

I-type

- ◆ Semantics

- $effective_address = sign_extend(offset) + GPR[base]$
- $MEM[translate(effective_address)] \leftarrow GPR[rt]$
- $PC \leftarrow PC + 4$


- ◆ Exceptions

- address must be “word-aligned”
- MMU exceptions



SW

Store Word



Format: SW rt, offset(base)

MIPS I

Purpose: To store a word to memory.

Description: $\text{memory}[\text{base} + \text{offset}] \leftarrow \text{rt}$

The least-significant 32-bit word of register *rt* is stored in memory at the location specified by the aligned effective address. The 16-bit signed *offset* is added to the contents of GPR *base* to form the effective address.

Restrictions:

The effective address must be naturally aligned. If either of the two least-significant bits of the address are non-zero, an Address Error exception occurs.

MIPS IV: The low-order 2 bits of the *offset* field must be zero. If they are not, the result of the instruction is undefined.

Operation: 32-bit Processors

```

vAddr ← sign_extend(offset) + GPR[base]
if (vAddr1..0) ≠ 02 then SignalException(AddressError) endif
(pAddr, uncached) ← AddressTranslation(vAddr, DATA, STORE)
dataword ← GPR[rt]
StoreMemory(uncached, WORD, dataword, pAddr, vAddr, DATA)

```

Operation: 64-bit Processors

```

vAddr ← sign_extend(offset) + GPR[base]
if (vAddr1..0) ≠ 02 then SignalException(AddressError) endif
(pAddr, uncached) ← AddressTranslation(vAddr, DATA, STORE)
pAddr ← pAddrPSIZE-1..3 || (pAddr2..0 xor (ReverseEndian || 02))
byte ← vAddr2..0 xor (BigEndianCPU || 02)
datadouble ← GPR[rt]63-8*byte || 08*byte
StoreMemory(uncached, WORD, datadouble, pAddr, vAddr, DATA)

```

Exceptions:

- TLB Refill, TLB Invalid
- TLB Modified
- Address Error



Assembly Programming 201

◆ E.g. High-level Code

```
A[ 8 ] = h + A[ 0 ]
```

where **A** is an array of integers (4-byte each)


◆ Assembly Code

- suppose **&A**, **h** are in r_A , r_h
- suppose r_{temp} is a free register

```
LW  $r_{temp}$  0( $r_A$ )      #  $r_{temp}$  = A[0]
add  $r_{temp}$   $r_h$   $r_{temp}$  #  $r_{temp}$  = h + A[0]
SW  $r_{temp}$  32( $r_A$ )     # A[8] =  $r_{temp}$ 
                          # note A[8] is 32 bytes
                          #      from A[0]
```



Load Delay Slots (old days..)



- ◆ R2000 load has an architectural latency of 1 inst*.
 - The instruction immediately following a load (in the “delay slot”) still sees the old register value
 - The load instruction no longer has an atomic semantics

Why would you do it this way?

- ◆ Is this a good idea?
(hint: R4000 redefined LW to complete atomically)

*BTW, notice the latency is defined in “instructions” not cyc. or sec.

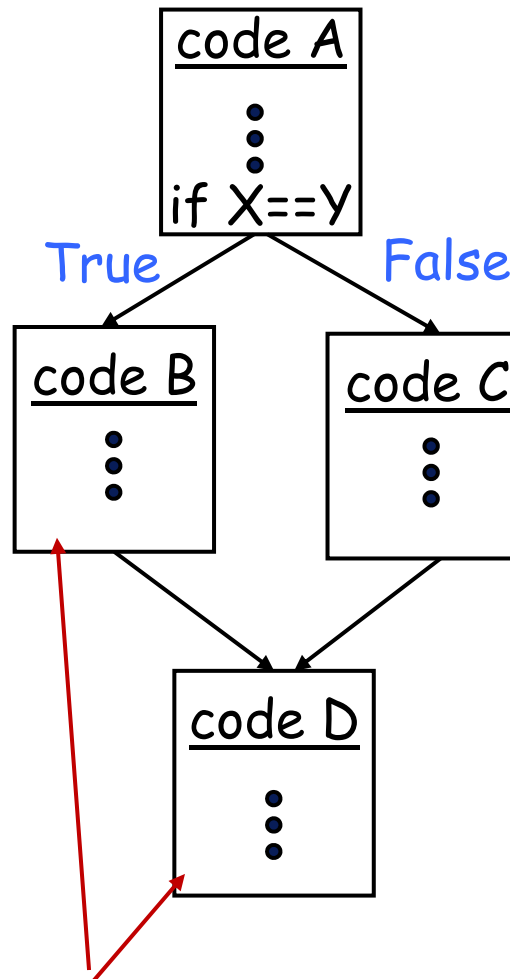


Control Flow Instructions

◆ C-Code


```
{ code A }  
if X==Y then  
    { code B }  
else  
    { code C }  
{ code D }
```

Control Flow Graph



these things are called **basic blocks**

Assembly Code (linearized)





(Conditional) Branch Instructions

- ◆ Assembly (e.g., branch if equal)

BEQ rs_{reg} rt_{reg} $immediate_{16}$

- ◆ Machine encoding

BEQ	rs	rt	immediate
6-bit	5-bit	5-bit	16-bit

I-type

- ◆ Semantics

// PC-relative jump

- $target = PC + \text{sign-extend}(immediate) \times 4 \quad \leftarrow (x\ 4) = (<<\ 2b)$
- if $GPR[rs] == GPR[rt]$ then $PC \leftarrow target$
else $PC \leftarrow PC + 4$

- ◆ How far can you jump?

- ◆ Variations


- BEQ, BNE, BLEZ, BGTZ

Why isn't there a BLT or BGT instruction?



BEQ

Branch on Equal



Format: BEQ rs, rt, offset

MIPS I

Purpose: To compare GPRs then do a PC-relative conditional branch.

Description: if (rs = rt) then branch

An 18-bit signed offset (the 16-bit *offset* field shifted left 2 bits) is added to the address of the instruction following the branch (**not** the branch itself), in the branch delay slot, to form a PC-relative effective target address.

If the contents of GPR *rs* and GPR *rt* are equal, branch to the effective target address after the instruction in the delay slot is executed.

Restrictions:

None

Operation:

```
I: tgt_offset ← sign_extend(offset || 02)  
   condition ← (GPR[rs] = GPR[rt])  
I+1: if condition then  
      PC ← PC + tgt_offset  
   endif
```

Exceptions:

None

Programming Notes:

With the 18-bit signed instruction offset, the conditional branch range is ± 128 KBytes.
Use jump (J) or jump register (JR) instructions to branch to more distant addresses.




Jump Instructions

- ◆ Assembly

J immediate₂₆

- ◆ Machine encoding



J-type

- ◆ Semantics

// PC-relative jump

- target = PC[31:28]x2²⁸ |_{bitwise-or} zero-extend(immediate) x 4
- PC ← target

- ◆ How far can you jump (in both direction)?

- ◆ Variations

- Jump and Link
- Jump Registers



Assembly Programming 301


◆ E.g. High-level Code

```
if (i == j) then
    e = g
else
    e = h
f = e
```

◆ Assembly Code

- suppose e, f, g, h, i, j are in $r_e, r_f, r_g, r_h, r_i, r_j$

```
    bne ri rj L1      # L1 and L2 are addr labels
                        # assembler computes offset
    add re rg r0      # e = g
    j L2
L1:  add re rh r0      # e = h
L2:  add rf re r0      # f = e
    . . . .
```





Branch Delay Slots

- ◆ R2000 branch instructions also have an architectural latency of 1 instructions
 - the instruction immediately after a branch is always executed (in fact PC-offset is computed from the delay slot instruction)
 - branch target takes effect on the 2nd instruction

```
bne ri rj L1  
  
add re rg r0  
j L2  
  
L1: add re rh r0  
  
L2: add rf re r0  
    . . . .
```



```
bne ri rj L1  
nop  
  
j L2  
add re rg r0  
L1: add re rh r0  
  
L2: add rf re r0  
    . . . .
```



Strangeness in the Semantics

Where do you think you will end up?

```
_s:  j L1
      j L2
      j L3

L1:  j L4
L2:  j L5

L3:  foo
L4:  bar
L5:  baz
```



Function Call and Return

- ◆ Jump and Link: JAL offset₂₆
 - return address = PC + 8
 - target = PC[31:28]x2²⁸ |_{bitwise-or} zero-extend(immediate)x4
 - PC ← target
 - GPR[r31] ← return address

On a function call, the callee needs to know where to go back to afterwards

- ◆ Jump Indirect: JR rs_{reg}
 - target = GPR [rs]
 - PC ← target


PC-offset jumps and branches always jump to the same target every time the same instruction is executed

Jump Indirect allows the same instruction to jump to any location specified by rs (usually r31)



Jump And Link

JAL



Format: JAL target

MIPS I

Purpose: To procedure call within the current 256 MB aligned region.

Description:

Place the return address link in GPR 31. The return link is the address of the second instruction following the branch, where execution would continue after a procedure call.

This is a PC-region branch (not PC-relative); the effective target address is in the “current” 256 MB aligned region. The low 28 bits of the target address is the *instr_index* field shifted left 2 bits. The remaining upper bits are the corresponding bits of the address of the instruction in the delay slot (not the branch itself).

Jump to the effective target address. Execute the instruction following the jump, in the branch delay slot, before jumping.

Restrictions:

None

Operation:

I: GPR[31] \leftarrow PC + 8
I+1: PC \leftarrow PC_{GPREN..28} || *instr_index* || 0²

Exceptions:

None

Programming Notes:

Forming the branch target address by catenating PC and index bits rather than adding a signed offset to the PC is an advantage if all program code addresses fit into a 256 MB region aligned on a 256 MB boundary. It allows a branch to anywhere in the region from anywhere in the region which a signed relative offset would not allow.

This definition creates the boundary case where the branch instruction is in the last word of a 256 MB region and can therefore only branch to the following 256 MB region containing the branch delay slot.



Assembly Programming 401

Caller

```
... code A ...  
JAL _myfxn  
... code C ...  
JAL _myfxn  
... code D ...
```

Callee

```
_myfxn:    ... code B ...  
          JR r31
```

- ◆ **A** \rightarrow call **B** \rightarrow return **C** \rightarrow call **B** \rightarrow return **D**
- ◆ How do you pass argument between caller and callee?
- ◆ If **A** set **r10** to 1, what is the value of **r10** when **B** returns to **C**?
- ◆ What registers can **B** use?
- ◆ What happens to **r31** if **B** calls another function



Caller and Callee Saved Registers

◆ Callee-Saved Registers

- Caller says to callee, “The values of these registers should not change when you return to me.”
 - Callee says, “If I need to use these registers, I promise to save the old values to memory first and restore them before I return to you.”
- “hey caller! I am giving you back original values on return.”

◆ Caller-Saved Registers

- Caller says to callee, “If there is anything I care about in these registers, I already saved it myself.”
 - Callee says to caller, “Don’t count on them staying the same values after I am done.”
- “hey callee! I saved it in my space. use it as you want.”

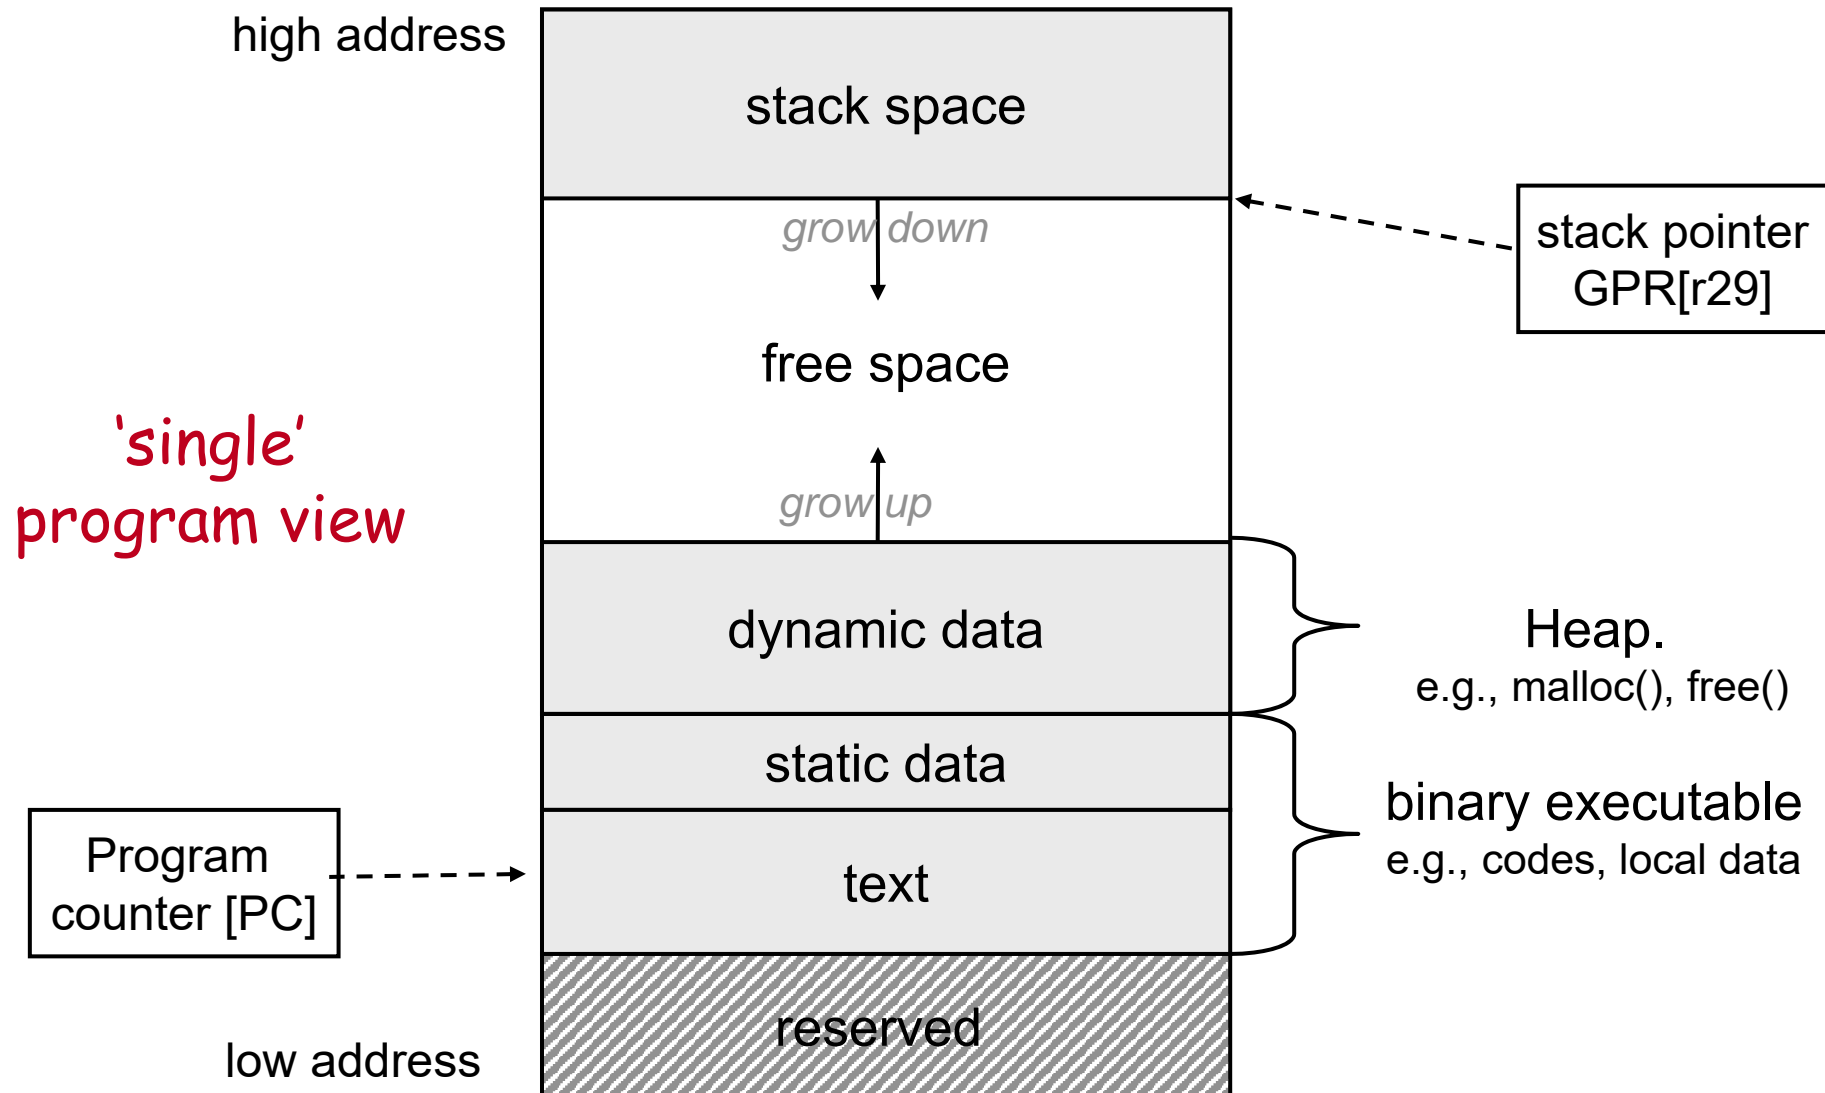


R2000 Register Usage Convention

- ◆ r0: always 0 = \$zero
- ◆ r1: reserved for the assembler = \$at
- ◆ r2, r3: function return values = \$v0 ~ \$v1
- ◆ r4~r7: function call arguments = \$a0 ~ \$a3
- ◆ r8~r15: “caller-saved” temporaries = \$t0 ~ \$t7
- ◆ r16~r23: “callee-saved” temporaries = \$s0 ~ \$s7
- ◆ r24~r25: “caller-saved” temporaries = \$t8 ~ \$t9
- ◆ r26, r27: reserved for the operating system = \$k0 ~ \$k1
- ◆ r28: global pointer = \$gp
- ◆ r29: stack pointer = \$sp
- ◆ r30: callee-saved temporaries = \$fp
- ◆ r31: return address = \$ra



R2000 Memory Usage Convention





Calling Convention

.....

1. caller saves caller-saved registers
2. caller loads arguments into r4~r7
3. caller jumps to callee using JAL

4. callee allocates space on the stack (dec. stack pointer)
5. callee saves callee-saved registers to stack (also r4~r7, old r29, **r31 <- why?**)

} prologue

..... body of callee (can "nest" additional calls)

6. callee loads results to r2, r3
7. callee restores callee-saved registers
8. JR r31

} epilogue

9. caller restores caller-saved registers
10. caller continues with return values in r2, r3



To Summarize: MIPS RISC

- ◆ Simple operations
 - 2-input, 1-output arithmetic and logical operations
 - few alternatives for accomplishing the same thing
- ◆ Simple data movements
 - ALU ops are register-to-register (need a large register file)
 - “Load-store” architecture
- ◆ Simple branches
 - limited varieties of branch conditions and targets
- ◆ Simple instruction encoding
 - all instructions encoded in the same number of bits
 - only a few formats

Mostly, ISA intended for compilers rather than assembly programmers.



We didn't talk about

- ◆ Privileged Modes
 - User vs. supervisor
- ◆ Exception Handling
 - trap to supervisor handling routine and back
- ◆ Virtual Memory
 - Each user has 4-GBytes of private, large, linear and fast memory?
- ◆ Floating-Point Instructions



Question?

Announcements: HW #1 will be posted on the course web site soon.

Reading: start reading P&H Ch.3 (stuff for the next class).

Handouts: None