

# Lecture 11: Exceptions and Interrupts

Jangwoo Kim (Seoul National University)

[jangwoo@snu.ac.kr](mailto:jangwoo@snu.ac.kr)

# Announcement

- ◆ HW#3
  - Posted!
  - **Due: 4/19**
  
- ◆ Mid-term
  - **4/20 (6PM) @ Room 102**
  
- ◆ One more lecture before mid-term exam
  - **Pick between 4/14 (2pm, Saturday) or 4/17 (5pm, Tuesday)**

**This will help you better prepare for mid-term exam**

# CPU must prepare for ‘unplanned’ events


- ◆ A systematic way to handle exceptional conditions that are relatively rare, but must be detected and acted upon quickly
  - Instructions may fail and cannot complete
  - External I/O devices may need servicing
  - Quantum expiration in a time-shared system
- ◆ Option 1: write every program with continuous checks (a.k.a. polling) for every possible contingency

acceptable for simple embedded systems (e.g., toaster)
- ◆ Option 2: write “normal” programs for the best-case scenario where nothing unusual happens
  - But, must detect exceptional conditions in HW
  - “Transparently” transfer control to an exception handler that knows how to resolve the condition and then back to your program

# Interrupt Control Transfer

- ◆ An interrupt is an “unplanned” function call to a system routine (a.k.a., the interrupt handler)
- ◆ Unlike a normal function call, the interrupted thread cannot anticipate the control transfer or prepare for it in any way
- ◆ Control is later returned to the main thread at the interrupted instruction

The control transfer to the interrupt handler and back must be 100% transparent to the interrupted thread!!!



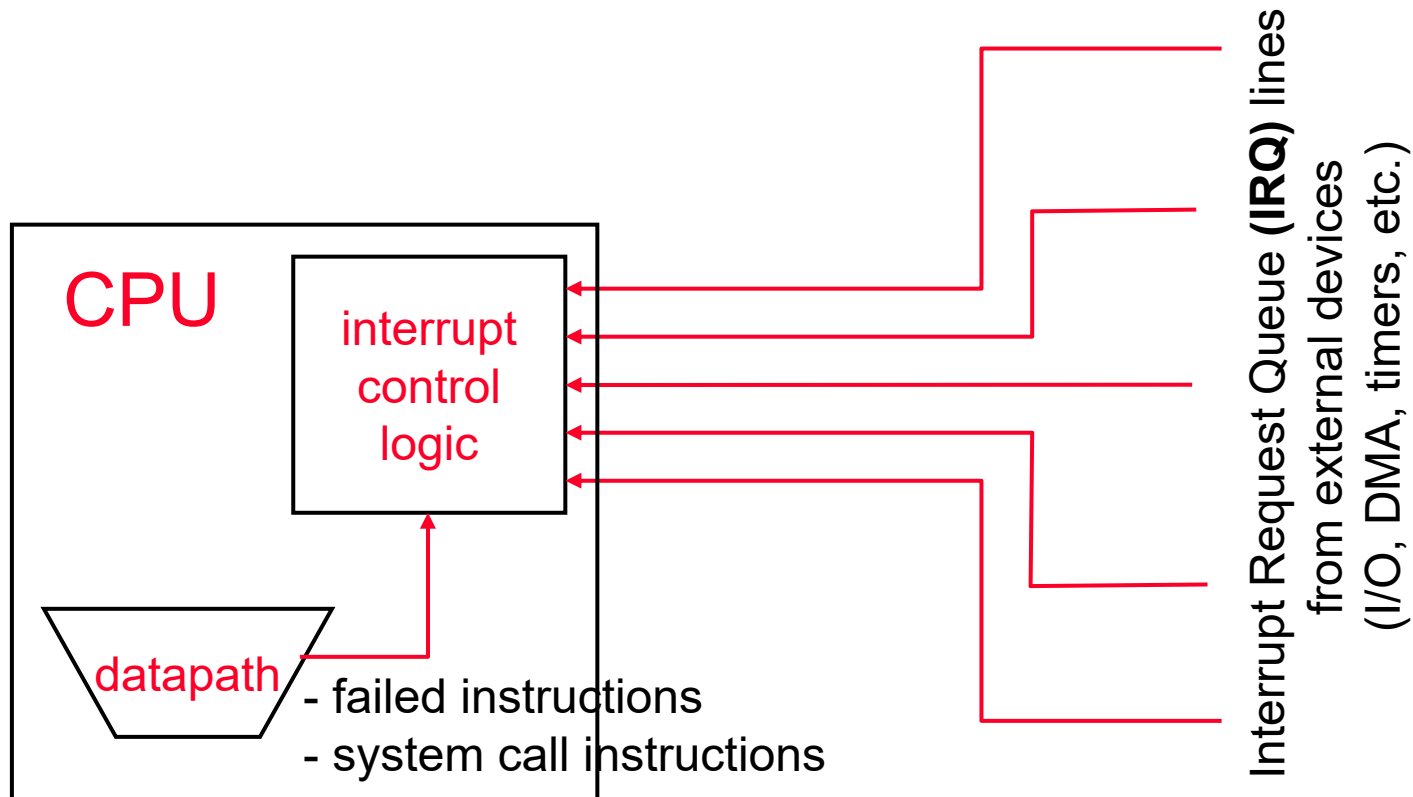
# Types of Interrupts

- ◆ Synchronous Interrupts (a.k.a., **exceptions**)
  - Exceptional conditions **tied to a particular instruction**  
e.g. illegal opcode, illegal operand, virtual memory management faults
  - The faulting instruction cannot be finished

**No forward-progress unless handled immediately**
- ◆ Asynchronous Interrupts (a.k.a., **interrupts**)
  - External events **not tied to a particular instruction**  
e.g., I/O events, timer events
  - Some flexibility on when to handle it

**Some delay allowed, but still cannot postpone forever**
- ◆ System Call/Trap Instruction
  - An instruction whose only purpose is **to raise an exception**

# Interrupt Sources



**Exceptions**  
(a.k.a. synchronous interrupts)

**Interrupts**  
(a.k.a. asynchronous interrupts,  
or external interrupts)

# Virtualization and Protection

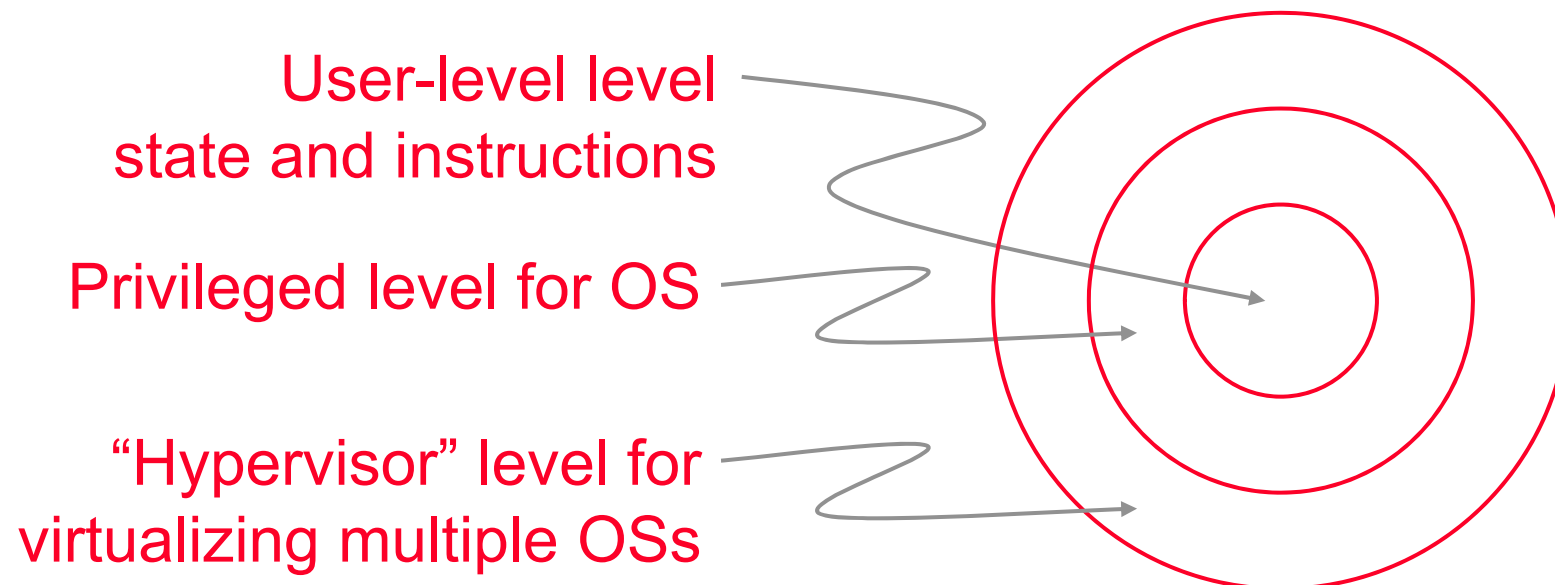
- ◆ Modern OS supports time-shared multiprocessing but each “user-level” process still thinks it is alone
  - Each process sees a private set of user-level architectural states that can be modified by the user-level instruction set
  - Each process cannot see or manipulate (directly) state and devices outside of this abstraction
- ◆ OS implements and manages a critical set of functionality
  - Keep low-level details out of the user-level process  
(e.g., HW device handler, thread scheduler, ..)
  - Protect the user-level process from each other and itself

Do you want to access/manage hard disk directly?

Do you trust your friend if he accesses the hard disk directly?

# Privilege Levels

- ◆ The OS must somehow be more powerful to create and maintain such an abstraction, hence a separate privileged (a.k.a. protected or kernel) mode
  - Additional architectural states and instructions, in particular to control virtualization/protection/isolation
  - The kernel code running in the privileged mode has access to the complete “bare” hardware system








# Control and Privilege Transfer

- ◆ User-level code never runs in the privileged mode
- ◆ Processor enters the privileged mode only **on interrupts**---user code surrenders control to a handler in the OS kernel
- ◆ The handler restores privilege level back to user mode before returning control to the user code




# Implementing Interrupts

# “Precise” Interrupt/Exception

Sequential Code Semantics

Overlapped Execution



A **precise interrupt** appears (to the interrupt handler) to take place exactly between two instructions

- Older instructions finished completely
- Younger instructions as if never happened
- On synchronous interrupts, execution stops just before the faulting instruction, and resume after interrupt handling done

# Stopping and Restarting a Pipeline

	$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$	$t_8$	$t_9$	$t_{10}$
IF	$I_0$	$I_1$	$I_2$	$I_3$	$I_4$	bub	bub	$I_h$	$I_{h+1}$	$I_{h+2}$	$I_{h+3}$
ID		$I_0$	$I_1$	$I_2$	$I_3$	bub	bub	bub	$I_h$	$I_{h+1}$	$I_{h+2}$
EX			$I_0$	$I_1$	$I_2$	bub	bub	bub	bub	$I_h$	$I_{h+1}$
MEM				$I_0$	$I_1$	$I_2$	bub	bub	bub	bub	$I_h$
WB					$I_0$	$I_1$	$I_2$	bub	bub	bub	bub

What if  $I_0$ ,  $I_1$ ,  $I_2$ ,  $I_3$  and  $I_4$  all generate exceptions in  $t_4$ ?  
How would things look different for asynchronous interrupts?

# Exception Sources in Different Stages

## ◆ IF

- Instruction memory address/protection fault

## ◆ ID

- Illegal opcode
- Trap to SW emulation of unimplemented instructions
- System call instruction (an intended exception requested by SW)

## ◆ EX

- Invalid results: overflow, divide-by-zero, ...

## ◆ MEM

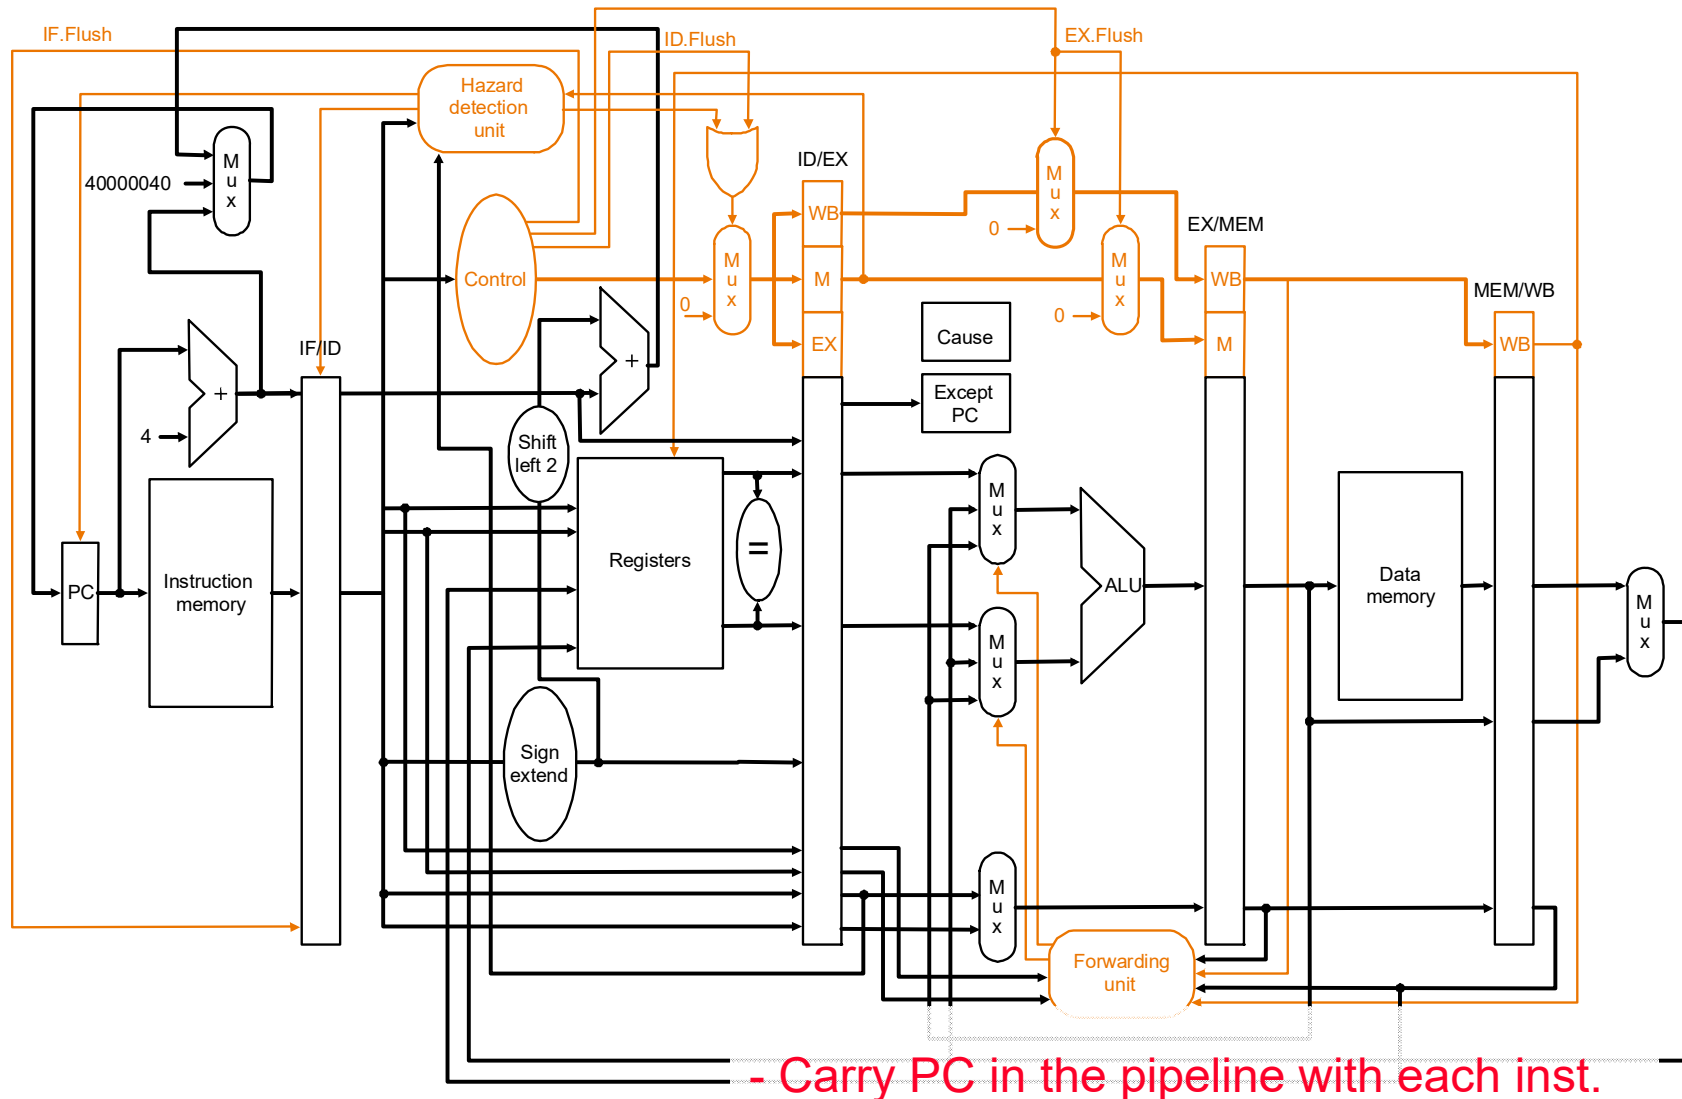
- Data memory address, protection fault

## ◆ WB

- Nothing can stop an instruction now...

- ◆ We can associate async interrupts (I/O) with any instruction/stage we like

# Pipeline Flush for Exceptions




- Carry PC in the pipeline with each inst.
- New pipeline flush points

# MIPS Interrupt Architecture


# MIPS Interrupt Architecture

## ◆ Privileged system control registers

- Exception Program Counter (EPC, CR14)
  - Which instruction did we stop on?
- Interrupt Cause Register (CR13) : What caused the interrupt?



- Interrupt Status Register (CR12): Enable and disable interrupts  
Set privilege modes



- ◆ Automatically set on interrupt transfer events
- ◆ Also accessed by the “move from/to control register-x” instruction: “mfc0 R<sub>y</sub>, CR<sub>x</sub>” and “mtc0 R<sub>y</sub>, CR<sub>x</sub>”



# MIPS Interrupt Architecture

- ◆ On an interrupt transfer, the CPU hardware saves the interrupt address to EPC
  - Can't just leave frozen in the PC: **overwritten immediately**
  - Can't use r31 as in a function call: **need to save user value**
- ◆ In general, CPU hardware must save any such information that cannot be saved and restored in software by the interrupt handler (**very few such things**)
- ◆ For example, the GPR can be managed in SW by the interrupt handler using a **callee-saved** convention
  - However, r26 and r27 are reserved by convention to be available to the kernel immediately at the start and the end of an interrupt handler
    - r26 & r27 (or \$k0~\$k1: reserved for OS kernel)

# Interrupt Servicing

- ◆ On an interrupt transfer, the CPU hardware records the cause of the interrupt in a privileged registers (Interrupt Cause Register, CR13)
- ◆ Option 1: Control is transferred to a pre-fixed default interrupt handler address **// flexible, but slow**
  - This initial handler examines the cause and branches to the appropriate handler subroutine to do the work
  - This address is protected from user-level process so one cannot just jump or branch to it
- ◆ Option 2: Vectored Interrupt **// not flexible, but fast**
  - A bank of privileged registers to hold a separate specialized handler address for each interrupt source
  - On an interrupt, hardware transfer control directly to the appropriate handler to save interrupt overhead

# Example of Causes

Number	Name	Cause of exception
0	Int	interrupt (hardware)
4	AdEL	address error exception (load or instruction fetch)
5	AdES	address error exception (store)
6	IBE	bus error on instruction fetch
7	DBE	bus error on data load or store
8	Sys	syscall exception
9	Bp	breakpoint exception
10	Rl	reserved instruction exception
11	CpU	coprocessor unimplemented
12	Ov	arithmetic overflow exception
13	Tr	trap
15	FPE	floating point

Figure from [P&H CO&D, COPYRIGHT 2004 Elsevier. ALL RIGHTS RESERVED.]

# Handler Examples

- ◆ On asynchronous interrupts, device-specific handlers are invoked to service the I/O devices
- ◆ On exceptions, kernel handlers are invoked to either
  - Correct the faulting condition and continue the program (e.g., emulate the missing FP functionality, update virtual memory management), or
  - “signal” back to the user process if a user-level handler function is registered, or
  - Kill the process if the exception cannot be corrected
- ◆ “System call” is a special kind of function call from user process to kernel-level service routines

# Returning from Interrupt

- ◆ Software restores all architectural states saved at the start of the interrupt routine

## (1) MIPS32 uses a special jump instruction

### - ERET

- > To atomically restore the automatically saved CPU states
- > To atomically restore the privilege level
- > To atomically jump back to the interrupted address in EPC

## (2) MIPS R2000 uses a pair of instructions

(assume “mfc0 r26,epc” done)

jr r26      // jump to a copy of EPC in r26

rfe         // restore from exception mode

            // must be used in the delay slot!!

# Branch Delay Slot and RFE

- ◆ What if the faulting address is a branch delay slot?
  - Simply jumping back to the faulting address won't continue correctly if the preceding branch was taken
  - We didn't save enough information to do the right thing
- ◆ MIPS's solution
  - CPU HW makes a note (in the Cause register) if the faulting address captured is in a delay slot
  - In these cases, the handler returns to the preceding branch instruction which gets executed twice
- ◆ Generally harmless except "JALR r31"
  - Explicitly disallowed by the MIPS ISA

What would happen if it is allowed?

# An Extremely Short Handler

`_handler_shortest:`

`# no prologue needed`

```
... short handler body ...  
# can use only r26 and r27  
# interrupt not re-enabled for  
# something really quick
```

`# epilogue`

`mfc0 r26,epc`

`jr r26`

`rfe`

`# get faulting PC`

`# jump to retry faulting PC`

`# restore from exception mode`

`= restore arch. state`

# A Short Handler

`_handler_short:`

`# prologue`

`addi sp, sp, -0x8`

`# allocate stack space (8 byte)`

`sw r8, 0x0(sp)`

`# back-up r8 and r9 for use in body`

`sw r9, 0x4(sp)`

`#`

`... short handler body ...`

`# can use r26, r27, and r8, r9`

`# interrupt not re-enabled`

`# epilogue`

`lw r8, 0x0(sp)`

`# restore r8, r9`

`lw r9, 0x4(sp)`

`#`

`addi sp, sp, 0x8`

`# restore stack pointer`

`mfc0 r26, epc`

`# get EPC`

`j r26`

`# jump to retry EPC`

`rfe`

`# restore from exception mode`



# Nesting Interrupts

- ◆ On an interrupt control transfer, further exceptions or asynchronous interrupts are disabled automatically
  - Another interrupt would overwrite the contents of the EPC and Interrupt Cause and Status Registers
    - (1) The handler must be carefully written not to generate synchronous exceptions itself during this window of vulnerability
    - (2) The handler must disable further asynchronous interrupts using interrupt status register (CR12)
- ◆ However, for long-running handlers, interrupt must be re-enabled not to miss critical interrupts
  - The handler must save the contents of EPC/Cause/Status to memory (stack) before re-enabling asynchronous interrupt
  - Once interrupts are re-enabled, EPC/Cause/Status can be updated by the next interrupt

# Interrupt Priority

- ◆ Asynchronous interrupt sources are ordered by priorities
  - Higher-priorities interrupts are more timing critical
  - If multiple interrupts are triggered, the handler handles the highest-priority interrupt first
- ◆ Interrupts from different priorities can be selectively disabled by setting the mask in the Status register
- ◆ When servicing a particular priority interrupt, the handler only re-enable higher-priority interrupts
  - Higher-priority interrupt should not get delayed

Re-enabling same/lower-priority interrupts may lead to an infinite loop if a device interrupts repeatedly

# Nestable Handler

\_handler\_nest:

# prologue

addi sp, sp, -0x8

mfc0 r26, epc

sw r26, 0x0(sp)

sw r8, 0x4(sp)

addi r26, r0, 0x405

mtc0 r26, status

# allocate stack space for EPC

# get EPC

# store EPC onto stack

# allocate a register for use later

# set interrupt enable bit

# write into status reg

... interruptible

# could free-up more registers

longer handler body ...

# to stack if needed

# epilogue

addi r8, r0, 0x404

mtc0 r8, status

ld r26, 0x0(sp)

ld r8, 0x4(sp)

addi sp, sp, 0x8

j r26

rfe

# clear interrupt enable bit

# write into status reg

# get EPC back from stack

# restore r8

# restore stack pointer

# jump to retry EPC

# restore from exception mode

# Question?

Announcements: Homework #3 will be collected (due: 4/19)

Mid-term on 4/20 (starting at 6:00PM)

Reading: Finish P&H Ch. 4 & Early part of Ch. 5

Handouts: None